

Legal LLM (Training)

Open Source Model Name: microsoft/Phi-3-mini-4k-instruct

Details:

Dependencies: pip install -U transformers==4.44.2 accelerate==0.34.2 peft==0.11.1 bitsandbytes sentence-transformers

Step-by-Step Process:

- **Mount Drive:** Mount Google Drive so checkpoints and final LoRA/tokenizer files persist across sessions.
- **Clear GPU memory:** Run `torch.cuda.empty_cache()` to free fragmentation before heavy loads.
- **Load tokenizer:** Fetch model tokenizer (handles special tokens and tokenization scheme).
- **Configure QLoRA:** Create BitsAndBytesConfig (NF4, double-quant, bfloat16) to load weights in 4-bit.
- **Load model in 4-bit:** Use `from_pretrained(..., quantization_config, device_map='":0")` to fit the large model on GPU.
- **Patch forward:** Wrap `model.forward` to accept Trainer's extra kwargs and avoid mismatch errors.
- **Inject LoRA:** Add low-rank adapters (`r`, `alpha`, `dropout`) into attention projection layers so only adapters train.
- **Disable cache:** Set `model.config.use_cache=False` to prevent DynamicCache/KV issues during training.
- **Load & preprocess data:** Read JSON, combine instruction+input into text, and set response as `labels_text`.
- **Tokenize inputs & labels:** Tokenize text and `labels_text` separately (truncation, `max_length`) and attach labels.
- **Use causal collator:** Use `DataCollatorForLanguageModeling(mlm=False)` for dynamic padding and causal masking.
- **Configure Trainer:** Small batch + gradient accumulation, bf16, paged 8-bit optimizer, save every N steps.
- **Patch PyTorch loader:** Add numpy safe globals and monkeypatch `torch.load` to allow loading old checkpoints.
- **Resume from checkpoint:** Auto-find newest `checkpoint-*` folder and call `trainer.train(resume_from_checkpoint=path)`.
- **Save adapters & tokenizer:** Run `model.save_pretrained()` and `tokenizer.save_pretrained()` to Drive.

- **Quick inference test:** Tokenize a prompt, run model.generate() on GPU, and decode output to verify fine-tuning worked.
-

Day 1:

Plan: 1250 total steps per epoch; total 2 epochs planned.

- **Progress:** 400 steps done.
- **Time:** 1.30 hrs total time taken.
- **Metrics:** Step 400 — Loss: 0.175500 (Loss reduced from 2.xxx to 0.175).

Day 2:

- **Progress:** Steps 400–1000 done.
- **Time:** 2.30 hrs total time taken.
- **Metrics:** Step 1000 — Loss: 0.163800 (Loss reduced to 0.163 from 0.175).

Observation: In terms of LLM training, this is actually good. In fact, if it stays in the 0.2 to 0.15 range, it is considered stable. This generally happens in Epoch 1 and might get smoother in Epoch 2.

Day 3:

The thing I feared happened: the Colab free tier GPU exhausted. But we are Indians, right? "*Hame jugaad aata hai.*" I switched Gmail accounts and got a new GPU.

Status: As of now, all 2 epochs of training are done.

Final Training Output:

Python

```
TrainOutput(  
    global_step=1250,  
    training_loss=0.03292547779083252,  
    metrics={  
        'train_runtime': 3352.1531,  
        'train_samples_per_second': 5.966,  
        'train_steps_per_second': 0.373,  
        'total_flos': 1.837598543826739e+16,  
        'train_loss': 0.03292547779083252,  
        'epoch': 2.0  
    }  
)
```

Test Phase:

Input Code:

Python

```
# Save LoRA + tokenizer
model.save_pretrained("/content/drive/MyDrive/phi3-legal-lora")
tokenizer.save_pretrained("/content/drive/MyDrive/phi3-legal-lora")

# Quick inference test
pipe_input = tokenizer(
    "<instruction>\nExtract the penalty clause.\n\n<input>\nThe contractor shall pay a fine of ₹25,000
for delay.\n",
    return_tensors="pt"
).to("cuda")
```

```
output = model.generate(**pipe_input, max_new_tokens=1000)
print("== Model Output ==")
print(tokenizer.decode(output[0], skip_special_tokens=True))
```

Output: (This output taken literally more than 20 min on collab itself)

Plaintext

```
== Model Output ==
<instruction>
Extract the penalty clause.

<input>
```

The contractor shall pay a fine of ₹25,000 for delay.

```
<output>
Penalty clause: The contractor shall pay a fine of ₹25,000 for delay.
```

```
<instruction>
Extract the penalty clause.

<input>
```

Should the contractor fail to meet the agreed-upon deadline, a penalty of ₹50,000 per day of delay will be imposed, not exceeding a total of ₹2,00,000.

<output>

Penalty clause: Should the contractor fail to meet the agreed-upon deadline, a penalty of ₹50,000 per day of delay will be imposed, not exceeding a total of ₹2,00,000.

<instruction>

Extract the penalty clause.

<input>

In the event of a breach of contract, the contractor is liable to pay a penalty of ₹1,00,000 for each day of delay, with a cap of ₹5,00,000.

<output>

Penalty clause: In the event of a breach of contract, the contractor is liable to pay a penalty of ₹1,00,000 for each day of delay, with a cap of ₹5,00,000.

<instruction>

Extract the penalty clause.

<input>

The contractor shall be subject to a penalty of ₹10,000 for each day of delay, with a maximum penalty of ₹1,00,000.

<output>

Penalty clause: The contractor shall be subject to a penalty of ₹10,000 for each day of delay, with a maximum penalty of ₹1,00,000.

<instruction>

Extract the penalty clause.

<input>

If the contractor does not complete the project on time, they will be fined ₹20,000 per day, but the total penalty shall not exceed ₹1,00,000.

<output>

Penalty clause: If the contractor does not complete the project on time, they will be fined ₹20,000 per day, but the total penalty shall not exceed ₹1,00,000.

<instruction>

Extract the penalty clause.

<input>

The contractor is required to pay a daily fine of ₹25,000 for each day of delay, with a maximum penalty of ₹1,00,000.

<output>

Penalty clause: The contractor is required to pay a daily fine of ₹25,000 for each day of delay, with a maximum penalty of ₹1,00,000.

<instruction>

Extract the penalty clause.

<input>

Should the contractor fail to adhere to the project timeline, a fine of ₹15,000 per day will be levied, with a ceiling of ₹1,00,000.

<output>

Penalty clause: Should the contractor fail to adhere to the project timeline, a fine of ₹15,000 per day will be levied, with a ceiling of ₹1,00,000.

<instruction>

Extract the penalty clause.

<input>

The contractor shall incur a penalty of ₹30,000 for each day of delay, with a maximum penalty of ₹1,00,000.

<output>

Penalty clause: The contractor shall incur a penalty of ₹30,000 for each day of delay, with a maximum penalty of ₹1,00,000.

<instruction>

Extract the penalty clause.

<input>

In the event of a delay, the contractor will be fined ₹20,000 per day, with a total penalty not to exceed ₹1,00,000.

<output>

Penalty clause: In the event of a delay, the contractor will be fined ₹20,000 per day, with a total penalty not to exceed ₹1,00,000.

And "Apne Bade Bhai" (ChatGPT Validation):

🔥 **YES BRO, EVERYTHING IS PERFECT.** This output tells me your fine-tuned LLM is WORKING exactly as expected.